A SMALL BUILD

AI

The AI words I was nodding along to.

Ontology. Knowledge graph. Semantic layer. Jev. I built a tiny tool to understand the difference.

A portfolio project on meaning, governance, and shared data.

Swipe →

THE PROBLEM

Every team brings AI ideas. Almost nobody can triage them.

1

Divergent vocabulary

Customer, risk, and value mean five different things.

2

Governance last

Controls arrive after the build, if at all.

3

Quick wins buried

High value and high risk sit side by side.

4

No shared foundation

Every use case rebuilds the same data.

THE INSIGHT

A score is an opinion. A gate is a rule.

SCORE ALONE

- A number invites arguing.
- High risk can be scored upward.
- Nothing stops a bad idea from winning.

ONTOLOGY PLUS GATES

- Shared meaning first.
- Hard gates before any score.
- A missing control cannot be out-scored.

Put the gates first. Then score what survives.

THE ONTOLOGY

An ontology is a fight about words, written down so the code can referee.

Not a diagram. An agreement: what a use case is, what risk means, which data is allowed, and how they connect.

Use case has an owner

Risk tier 3 needs a human check

Approved data only

Customer links to account

Meaning, written down.

THE TOOL

One pipeline, driven by the ontology.

Validate

Fields and tiers



Gate

Owner, control, source



Score

0 to 100, explained



Bucket

Now, Next, Later, Hold



Shared layer

Entity reuse and cover

THE GATE

A high score cannot rescue a missing control.

Use case	Risk	Score	Outcome
Policy Q&A	1	98.8	Now
KYC Extraction	3	81.6	Now
Credit Memo	3	69.0	Hold

The Credit Memo scores 69 and still stops. Risk tier 3, no named owner, no monitoring, and an unapproved source.

The gate holds the line, not the number.

THE BUCKETS

Four buckets, four rules.

NOW

5

Score 70+, ready, low effort, no gaps.

NEXT

4

Score 55+, no gaps.

LATER

1

No gaps, low score.

HOLD

2

Any governance gap.

One foundation, not twelve.

Customer data showed up in 7 of 12 use cases.



Recommended layer: customer, policy, document, call_transcript.
100% of the sample can run on it.

HONESTY

What is real, what is invented, what is a judgement.

REAL

The ontology, the gates, the deterministic score, the shared-layer maths.

INVENTED

The 12 use cases and the 1 to 5 values. Synthetic, not employer data.

JUDGEMENT

The weights, the bucket rules, and the simplified risk tiers.

A portfolio tool only works if readers can tell the three apart.

The model is the easy part. The agreement is the work.

AI Use Case Portfolio Analyzer

github.com/noxnoxdata/ai-use-case-portfolio

Bridget Moh

Arguing with a dictionary